

Interrogating with AI

Exploring the Usage of and Auditing of Large Language Models in
Relational Coordination Research

Kartik Trivedi

2025-11-07

University of New Hampshire / RCC Roundtable 2025

Agenda

- What are LLMs?
- What changes in qualitative inquiry with GenAI?
- Opportunities + risks in QDA
- Prompt engineering & **parameter sensitivity**
- The **MERIT** framework for responsible reporting
- Hands-on with LLM tools
- Reflection & Q&A

What is a Large Language Model (LLM)?

- AI trained on **massive** text data
- Learns patterns → generates & interprets language
- Can summarize, code, classify, compare
- Examples: GPT-4, Claude, LLaMA

Why “Large” Matters

- Billions of parameters → emergent reasoning
- Transformer architecture (“attention”)
- Handles complexity & longer qualitative texts
- Limitations: hallucinations, bias, data opacity

Types of LLMs

Category	Examples	Why it matters
Model family	GPT, Claude	Capabilities differ
Openness	LLaMA vs GPT-4	Reproducibility
Modality	Text vs multimedia	Conversation types
Interface	App vs API vs QDA tools	Workflow impact

What is Qualitative Data Analysis?

- Interpretation, meaning, lived experience
- Coding [?] themes [?] narrative
- Always reflexive & relational

Traditional QDA Workflow

1 Data familiarization 2 Open/initial coding 3 Theme development 4 Interpretation & checking 5 Reporting + reflexivity

Why Use LLMs for QDA?

- ☐ First-pass coding + summarization
- ☐ Handle large corpora
- ☐ Cognitive support
- ☐ Speed for iteration
- ☐ Must maintain trustworthiness & participant dignity

Technologies Are Not Neutral

- Early view: digital tools as neutral *instruments*
- Updated view: tools shape inquiry
- LLMs have **power** in what is seen / unseen

Key Threats of GenAI

- ☐ Exploited data workers
- ☐ Environmental burden
- ☐ IP + consent issues
- ☐ Algorithmic bias
- ☐ Hallucination errors
- ☐ Loss of authenticity

➔ Engagement with AI is a **choice**, not a destiny

Efficiency ≠ Epistemic Quality

- “Faster” is not a valid analytic paradigm
- Summary ≠ interpretation
- LLMs flatten nuance if unchecked

Prompt Engineering = Analytic Intervention

How we ask affects:

- Codes generated
- Voices prioritized
- Power dynamics in text

Prompt Sensitivity Example

- **Prompt A** → high-level themes
- **Prompt B** → structured codebook with quotes

Differences are **methodological**, not cosmetic

LLM Parameters: What They Control

Parameter	Meaning	Impact on QDA
Temperature	Creativity vs precision	Nuance vs drift
Top-p / Top-k	Diversity of ideas	Breadth of coding
Max Tokens	Length/depth	Truncation vs saturation
Model choice	Training biases	Interpretive differences
Context window	Memory	Linking across data
System role	Interpretive lens	Analytical stance
Repetition penalty	Novelty	New vs redundant codes

Example: Changing Findings

□ Same transcript □ Same prompt □ Different parameters

Temp 0.1	Temp 0.9
Literal codes	Emotional interpretation
Conservative grouping	Diverse yet unstable themes
High repeat	Higher hallucination risk

➔ Meaning changes ➔ So must **report** settings

Parameters & Rigor

Use **MERIT** to evaluate effects:

M | Methods shift with settings |

E | Whose perspectives are emphasized/silenced? |

R | Who validates decisions? |

I | Did quality improve or degrade? |

T | Is the tool transparent about settings? |

Transparency Matters

- QDA already criticized for “black box” analysis
- GenAI increases opacity
- MERIT encourages explicit reporting:
 - Prompts
 - Model role & settings
 - Human oversight

Summary of MERIT Framework

- **M**ethods
- **E**thics
- **R**esponsibility
- **I**mpact
- **T**ool

□ A reflexive guide for **trustworthy** GenAI practices

Hands-On Exploration

You will: 1. Run a baseline prompt 2. Adjust parameters 3. Compare changes 4. Reflect using MERIT

Your Custom Tool Demo

- Modify: temp, top-p, tokens, model, role
- Compare outputs side-by-side
- Export prompt logs
- Built for **transparency in analytic decisions**

Best Practices

- ☐ Human-in-the-loop
- ☐ Prompt logs saved
- ☐ Model version + date recorded
- ☐ Validate outputs manually
- ☐ Member/peer checking
- ☐ Consent for AI use with participant data

Reflection Discussion

- ☐ What worked?
- ☐ What didn't feel trustworthy?
- ☐ Where will GenAI enter *your* workflow responsibly?

References (APA Suggested)

- Silver, C., Paulus, T., & ... (2024). *Evaluating GenAI in QDA...*
- Paulus, T., & Lester, J. (2021). *Doing Qualitative Research in a Digital World*.
- Minaee, S., et al. (2024). *Large Language Models: A Survey*.
- Zhu, et al. (2023). *LLM-in-the-loop Thematic Analysis*.
- Hannah & Bender (2022). *The myth of AI inevitability*.

Thank You

Questions & Discussion

Licensed under CC BY 4.0